

# **Architecture Document: Multi-Tenant SaaS on AWS (Commercial & GovCloud) with EKS and FIPS/DoD Posture via CloudFormation**

Author: Ted Gibson | Version: 2.7 | Date: 2025-08-28

## Table of Contents

- **Acronyms & Definitions**
- **Chapter 1 — Executive Summary**
- **Chapter 2 — SaaS Principles**
- **Chapter 3 — Infrastructure Layer**
  - 3.1 AWS Accounts, VPCs, Subnets, **Security Groups**, Endpoints
  - 3.2 EKS Clusters (dev/stg/prod)
  - 3.3 Add-ons (ALB Controller, Metrics Server, OIDC Integration)
  - 3.4 **Preflight Validation Scripts (deploy/preflight-vpc.sh)**
  - 3.5 **FIPS Endpoints, VPC-only, Compliance Controls**
- **Chapter 4 — CI/CD Pipeline Layer**
  - 4.1 Full CodePipeline Flow (Source → Build → Scan → Sign/Attest → Approval → Deploy → Mirror)
  - 4.2 Why CodePipeline + CodeBuild (FedRAMP Native)
  - 4.3 Tools Inside CodeBuild (Syft, Trivy, Cosign)
  - 4.4 FIPS Enforcement (VPC Endpoints)
  - 4.5 Repo Proof: Buildspecs + Pipeline Templates
- **Chapter 5 — Security & Compliance**
  - 5.1 Runtime Enforcement (Kyverno verifyImages, PodSecurity)
  - 5.2 IAM Least Privilege
  - 5.3 IRSA ServiceAccounts & OIDC Federation
  - 5.4 Zero Trust Posture
  - 5.5 Compliance Tie-In (FedRAMP, RMF, FIPS baked in)
  - 5.6 **Secrets Management (AWS Secrets Manager + External Secrets Operator)**
- **Chapter 6 — Supply Chain Security**
  - 6.1 SBOM Generation (Syft)
  - 6.2 Vulnerability Scanning (Trivy)
  - 6.3 Signing + Attestation with KMS (Cosign)
  - 6.4 Kyverno Admission Checks (verifyImages, PodSecurity)
  - 6.5 Immutable ECR Repositories
- **Chapter 7 — Tenant Lifecycle**
  - 7.1 Namespace-per-Tenant
  - 7.2 Helm Tenant Templates
  - 7.3 Quotas & Network Policies
  - 7.4 Onboarding Automation (bootstrap-fill-and-deploy.sh)
  - 7.5 Noisy-Neighbor Protection & Isolation Models
- **Chapter 8 — Public vs Private Partition (Commercial vs GovCloud)**
  - 8.1 Commercial → GovCloud Mirroring
  - 8.2 GovCloud Restrictions (Private-only, FIPS Endpoints)
  - 8.3 Why This Dual Setup is Necessary

- **Chapter 9 — Approval & Governance**
  - 9.1 Manual Approvals in CodePipeline (dev → stg → prod)
  - 9.2 Signed Artifact Verification
  - 9.3 Notifications & Event Routing (SNS, EventBridge, Slack/ServiceNow)
  - 9.4 Ops Handoff
- **Chapter 10 — Observability & Operations**
  - 10.1 Logging (CloudWatch, S3, Fluent Bit/ELK optional)
  - 10.2 Metrics (CW Container Insights, AMP, Grafana)
  - 10.3 Tracing (X-Ray / OpenTelemetry)
  - 10.4 Cost Allocation (CUR, Tags, Grafana overlays)
  - 10.5 **Backup, DR & Lifecycle Management**
  - 10.6 **Cost & Resilience Guardrails (ECR lifecycle, DR runbook, tenant tagging)**
- **Chapter 11 — Criteria Compliance Matrix**
- **Chapter 12 — Conclusion**

## Appendices

- Appendix A — AWS Services by Function
- Appendix B — Environment Setup & Validation Checklist
- Appendix C — Deployment Evidence (Scripts & Outputs)
- Appendix D — File-to-Chapter Mapping (Repo Proof)
- Appendix E — **Optional AWS Services for Scale-Out (Service Catalog, Control Tower)**

## Acronyms & Definitions

- **ALB** – Application Load Balancer
- **AMG** – Amazon Managed Grafana
- **AMP** – Amazon Managed Prometheus
- **API** – Application Programming Interface
- **ATO** – Authority to Operate
- **AWS** – Amazon Web Services
- **CAC** – Common Access Card
- **CI/CD** – Continuous Integration / Continuous Delivery
- **CJIS** – Criminal Justice Information Services (FBI) security policy
- **CM-8** – NIST 800-53 Control: Information System Component Inventory
- **CodeBuild** – AWS managed build service for CI/CD
- **CodeCommit** – AWS managed Git-based source control service
- **CodePipeline** – AWS managed CI/CD orchestration service
- **CUR** – Cost and Usage Report (AWS Billing)
- **Dev/Stg/Prod** – Development / Staging / Production environments
- **DoD SRG** – Department of Defense Security Requirements Guide
- **DR** – Disaster Recovery
- **EBS** – Elastic Block Store
- **ECR** – Elastic Container Registry

- **EFS** – Elastic File System
- **EKS** – Elastic Kubernetes Service
- **ELK** – Elasticsearch, Logstash, Kibana stack
- **EO 14028** – Executive Order on Improving the Nation's Cybersecurity (Supply Chain)
- **ESO** – External Secrets Operator
- **FedRAMP** – Federal Risk and Authorization Management Program
- **FIPS** – Federal Information Processing Standard (e.g., 140-2 for cryptography)
- **HPA** – Horizontal Pod Autoscaler
- **IAM** – Identity and Access Management
- **IRSA** – IAM Roles for Service Accounts (Kubernetes → IAM integration)
- **ITAR** – International Traffic in Arms Regulations
- **KMS** – Key Management Service (AWS cryptographic key management)
- **Kyverno** – Kubernetes-native policy engine (admission controller)
- **NIST 800-37 (RMF)** – Risk Management Framework
- **NIST 800-53** – Security and Privacy Controls for Information Systems
- **NLB** – Network Load Balancer
- **OIDC** – OpenID Connect (federated identity standard)
- **OTel** – OpenTelemetry (observability framework)
- **PIV** – Personal Identity Verification (government smart card)
- **RBAC** – Role-Based Access Control
- **RMF** – Risk Management Framework (NIST 800-37)
- **S3** – Simple Storage Service
- **SBOM** – Software Bill of Materials
- **SCP** – Service Control Policy
- **Seccomp** – Secure Computing Mode (Linux kernel sandboxing)
- **SES** – Simple Email Service
- **SNS** – Simple Notification Service
- **SQS** – Simple Queue Service
- **STS** – Security Token Service
- **Syft** – SBOM generation tool
- **Tenant Stamp** – A fully isolated environment (VPC, EKS, pipeline) provisioned per customer
- **Trivy** – Vulnerability scanning tool
- **VPC** – Virtual Private Cloud
- **WAF** – Web Application Firewall
- **Zero Trust** – Security framework requiring continuous verification

## Chapter 1 — Executive Summary

### 1.1 What the SaaS Is

This solution delivers a **multi-tenant Software-as-a-Service (SaaS) platform** designed to operate across both **AWS Commercial** and **AWS GovCloud** partitions. It provides a unified codebase with configurable deployment models:

- **Pooled:** Shared infrastructure (Aurora PostgreSQL with Row Level Security, shared EKS cluster namespaces).
- **Siloed:** Tenant-specific schemas and namespaces within shared clusters.
- **Dedicated Stamp:** Fully isolated infrastructure stack (separate AWS account, VPC, EKS, and CI/CD pipeline).

This architecture enables federal and commercial customers to adopt **zero-trust scoring and compliance coaching** services at the level of isolation their mission requires.

### 1.2 Why AWS

AWS was chosen because it uniquely provides:

- **FedRAMP High authorization** across Commercial and GovCloud services.
- **FIPS 140-2 validated endpoints** for cryptographic operations (KMS, S3, ECR, STS).
- **Scalable Kubernetes runtime (Amazon EKS)** with add-ons for ingress, autoscaling, and observability.
- **Mature security ecosystem** including IAM, VPC endpoints, WAF, and KMS.
- **Operational parity** between Commercial and GovCloud, enabling a mirrored pipeline without duplicate build systems.

### 1.3 Why Multi-Tenant

Delivering as multi-tenant SaaS provides:

- **Cost efficiency:** Shared infrastructure where possible (pooled/silo) with optional full isolation.
- **Flexibility:** Support for commercial, DoD, and federal agencies with one codebase.
- **Consistency:** Standardized CI/CD, governance, and compliance guardrails across all tenants.
- **Speed:** Tenant onboarding via Helm templates and automation scripts (bootstrap-fill-and-deploy.sh).

### 1.4 Why FIPS/GovCloud

Some customers mandate operations in **GovCloud** to satisfy **ITAR, CJIS, and DoD SRG** requirements. This platform supports:

- **GovCloud deployment parity** with Commercial builds (signed images mirrored via ECR).
- **FIPS 140-2 validation** enforced by VPC endpoints and private-only builds.
- **Controlled operations** managed by U.S. persons only in GovCloud.
- **Secure CI/CD handoff** between Commercial and GovCloud partitions without duplicating build pipelines.

### 1.5 Key Drivers: Compliance, Scalability, Security, Repeatability

- **Compliance:** FedRAMP High, FIPS 140-2, and NIST RMF (800-37/800-53) controls are baked into every layer.

- **Scalability:** EKS clusters with Horizontal Pod Autoscaler (HPA) and AWS Load Balancer Controller. Tenants can be deployed in pooled, siloed, or dedicated stamps with minimal configuration.
- **Security:** Zero Trust is enforced at build (SBOM + scans + cosign signatures), deploy (immutable digests), and runtime (Kyverno admission controls).
- **Repeatability:** All resources are provisioned via **Infrastructure-as-Code (CloudFormation + Helm)**, ensuring deterministic deployments and auditability.

## Chapter 2 — SaaS Principles

### 2.1 Single Codebase, Multiple Deployment Models

#### Understanding:

The repo is designed around a **single codebase** that supports three levels of tenant isolation: pooled, siloed, and dedicated stamps. This reduces complexity while offering flexibility for different compliance and mission requirements.

#### Approach:

- **Pooled Model:** Tenants share EKS clusters and Aurora PostgreSQL (Serverless v2) with **Row Level Security (RLS)**. Enforced via Cognito JWT claims and DB policies.
- **Silo Model:** Tenants share infrastructure but operate with **dedicated DB schemas, namespaces, IAM roles, and quotas**.
- **Dedicated Stamp Model:** A complete **isolated AWS account + VPC + EKS + pipeline** deployed per tenant, typically for federal or classified workloads.

#### Tools & Platforms:

- Kubernetes namespaces, Helm templating
- Aurora PostgreSQL with RLS policies (ops/migrations/aurora/0001\_init\_tenancy.sql)
- CloudFormation templates for isolated “stamps” (templates/aurora-postgres-slsv2.yaml)

#### Proof in Repo:

- ops/tenant/values-tenant-template.yaml
- ops/k8s/tenant-networkpolicy.yaml
- ops/k8s/tenant-resourcequota.yaml

#### Benefit:

Delivers **cost efficiency** for commercial tenants, while offering **mission-specific isolation** for government agencies — all without maintaining multiple codebases.

### 2.2 Public (Commercial) vs Private (GovCloud) Separation

#### Understanding:

Many federal agencies require workloads in GovCloud due to **ITAR, CJIS, and DoD SRG** compliance. The repo supports **pipeline parity** across Commercial and GovCloud to maintain consistency.

#### Approach:

- **Commercial Partition:** Builds, scans, SBOM generation, and artifact signing occur here using the full AWS service catalog.

- **GovCloud Partition:** Receives only **signed + attested artifacts** from Commercial via ECR mirroring (templates/gov-mirror-receiver.yaml).
- **No internet egress** allowed in GovCloud — all API traffic routed through **FIPS VPC endpoints**.

#### Tools & Platforms:

- AWS CodePipeline / CodeBuild
- AWS ECR (cross-partition replication)
- AWS KMS for signing (Cosign integration)

#### Proof in Repo:

- ops/pipeline/buildspec-mirror-gov.yaml
- templates/gov-mirror-receiver.yaml

#### Benefit:

Guarantees **audit parity** across partitions, while ensuring **federal workloads run inside GovCloud** without duplicating build pipelines.

## 2.3 Customer Isolation: Namespaces, IAM, Quotas, and Network Segmentation

#### Understanding:

Multi-tenancy requires strong isolation at runtime to prevent data leakage and ensure zero trust.

#### Approach:

- **Namespaces:** Each tenant runs in a dedicated Kubernetes namespace.
- **IAM Roles:** Workloads assume scoped roles via IRSA (IAM Roles for Service Accounts).
- **Resource Quotas:** Ensure no tenant can overconsume compute, memory, or storage.
- **Network Policies:** Restrict pod-to-pod and pod-to-service communications, preventing lateral movement.

#### Tools & Platforms:

- Kubernetes ResourceQuotas, NetworkPolicies, RBAC
- AWS IAM + IRSA integration (templates/workload-deploy-iam.yaml)

#### Proof in Repo:

- ops/k8s/tenant-networkpolicy.yaml
- ops/k8s/tenant-resourcequota.yaml

#### Benefit:

Delivers **noisy-neighbor protection**, **data isolation**, and **audit-ready compliance evidence** for FedRAMP and NIST RMF.

## 2.4 Zero Trust by Design

#### Understanding:

OMB M-22-09 and CISA Zero Trust Maturity Model mandate continuous verification at every layer. This repo enforces Zero Trust principles in **identity, network, runtime, and pipeline**.

#### Approach:

- **Identity:** Cognito/OIDC federation, IAM least privilege, IRSA.
- **Network:** Private-only VPC subnets, FIPS endpoints, and WAF.

- **Runtime:** Kyverno admission policies enforcing signed images + PodSecurity.
- **Pipeline:** Supply chain integrity enforced via SBOMs, cosign signatures, and manual approvals.

**Proof in Repo:**

- ops/policies/kyverno-verifyimages.yaml
- ops/policies/kyverno-podsecurity.yaml
- deploy/preflight-vpc.sh

**Benefit:**

Creates a **compliance-first SaaS platform** where every layer is secured, monitored, and auditable.

## Chapter 3 — Infrastructure Layer

### 3.1 AWS Accounts, VPCs, Subnets, Security Groups, Endpoints

**Understanding:**

A compliant SaaS must separate workloads across accounts and enforce **least-privilege networking**. The repo uses AWS Organizations and CloudFormation templates to provision accounts and VPCs for tooling (CI/CD), dev, stg, and prod.

**Approach:**

- **Accounts:** Tooling account for CI/CD, workload accounts for dev/stg/prod, optional dedicated accounts for stamps.
- **VPCs:** Dedicated VPCs per environment. Private subnets for workloads, public subnets only for load balancers.
- **Subnets:** Multi-AZ private subnets to ensure HA.
- **Security Groups:** Restrictive ingress/egress, allowing only required service flows (e.g., ALB → EKS, EKS → RDS).
- **Endpoints:** VPC Interface Endpoints enforce FIPS 140-2 service calls for S3, ECR, KMS, STS, CodeBuild, Logs.

**Tools & Platforms:**

- AWS Organizations, VPC, SGs, VPC Endpoints
- IaC via templates/\*.yaml, params/\*.json

**Proof in Repo:**

- deploy/preflight-vpc.sh (validation of endpoints)
- templates/tools-pipeline.yaml
- params/tools-\*.json

**Benefit:**

Network isolation, FIPS enforcement, and account separation ensure **Zero Trust**, prevent data exfiltration, and provide audit-ready evidence.



### 3.2 EKS Clusters (dev/stg/prod)

#### Understanding:

All workloads must move through gated environments with **parity across clusters** to ensure consistency.

#### Approach:

- **EKS Clusters:** One per environment (eks-dev, eks-stg, eks-prod).
- **Parity:** Configured identically using Helm and CloudFormation.
- **IAM Governance:** Deploy roles scoped per environment and mapped into EKS aws-auth ConfigMaps.

#### Tools & Platforms:

- Amazon EKS
- IAM + IRSA
- Helm

#### Proof in Repo:

- ops/helm/values-dev.yaml
- ops/helm/values-stg.yaml
- ops/helm/values-prod.yaml
- templates/workload-deploy-iam.yaml

#### Benefit:

Supports **promotion gates** (dev → stg → prod), ensures consistency, and enforces least privilege per environment.

### 3.3 Add-ons (ALB Controller, Metrics Server, OIDC Integration)

#### Understanding:

Certain add-ons are required for scaling, ingress, and secure identity integration.

#### Approach:

- **AWS Load Balancer Controller:** Provides ALBs for tenant ingress.
- **Metrics Server:** Required for HPA to scale workloads.
- **OIDC Integration:** Associates each EKS cluster with an OIDC provider, enabling IRSA.

#### Tools & Platforms:

- ALB Controller Helm chart
- Kubernetes Metrics Server
- AWS OIDC Provider

#### Proof in Repo:

- ops/helm/templates/\*
- Cluster OIDC setup documented in deploy/bootstrap-fill-and-deploy.sh

#### Benefit:

Ensures **elastic scale**, **secure ingress**, and **tenant-aware IAM integration**.

### 3.4 Preflight Validation Scripts (deploy/preflight-vpc.sh)

#### Understanding:

Before deploying workloads, the architecture must validate network and FIPS readiness.

#### Approach:

- Script checks for required VPC endpoints: ecr.api, ecr.dkr, s3, logs, sts, kms, codecommit, codebuild, events, ec2, eks.
- Fails early if any endpoint is missing.
- Enforces private-only builds (no NAT/public egress).

#### Tools & Platforms:

- Shell script in deploy/preflight-vpc.sh

#### Proof in Repo:

- deploy/preflight-vpc.sh

#### Benefit:

Prevents misconfigurations, enforces **Zero Trust networking**, and guarantees **FIPS-compliant API calls** before workloads ever launch.

### 3.5 FIPS Endpoints, VPC-only, Compliance Controls

#### Understanding:

Ongoing compliance requires **all cryptographic operations and service calls** to run through FIPS 140-2 validated endpoints.

#### Approach:

- VPC Endpoints configured with FIPS DNS suffixes.
- CodeBuild/CodePipeline run in private subnets, no internet egress.
- Security Groups and Network ACLs ensure least-privilege flows.

#### Tools & Platforms:

- AWS KMS, ECR, S3, STS with FIPS endpoints
- IaC templates + validation scripts

#### Proof in Repo:

- deploy/preflight-vpc.sh
- templates/tools-pipeline.yaml

#### Benefit:

Delivers **compliance-first assurance** (FedRAMP High, FIPS 140-2, NIST RMF), prevents audit failures, and reduces ATO timelines.

## Chapter 4 — CI/CD Pipeline Layer

### 4.1 Full CodePipeline Flow (Source → Build → Scan → Sign/Attest → Approval → Deploy → Mirror)

#### Understanding:

Section 10 of the architecture requires a **secure, auditable CI/CD pipeline** that enforces supply chain security, gated approvals, and parity between Commercial and GovCloud.

#### Approach:

- **Source:** CodeCommit repository serves as source of truth.
- **Build:** CodeBuild compiles workloads into container images.
- **Scan:** SBOMs generated (Syft), vulnerabilities scanned (Trivy).
- **Sign/Attest:** Cosign signs artifacts with AWS KMS asymmetric keys.
- **Approval:** Manual approvals inserted between dev → stg → prod.
- **Deploy:** Helm deploys immutable digests into EKS clusters.
- **Mirror:** Signed + attested artifacts mirrored to GovCloud ECR.

#### **Tools & Platforms:**

- AWS CodePipeline, CodeCommit, CodeBuild, KMS, ECR, Helm

#### **Proof in Repo:**

- templates/tools-pipeline.yaml
- ops/pipeline/buildspec-build.yml
- ops/pipeline/buildspec-deploy.yml
- ops/pipeline/buildspec-mirror-gov.yml

#### **Benefit:**

Provides an **end-to-end immutable chain of custody**, enforcing **compliance-first deployments**.

### **4.2 Why CodePipeline + CodeBuild (FedRAMP Native)**

#### **Understanding:**

Only AWS-native CI/CD services are FedRAMP High authorized in both Commercial and GovCloud.

#### **Approach:**

- **CodePipeline:** Orchestrates CI/CD flow, approvals, notifications.
- **CodeBuild:** Executes builds in private subnets, ensuring all service calls use **FIPS endpoints**.
- **Native Integration:** CloudWatch, IAM, KMS, and SNS integrated out-of-the-box.

#### **Proof in Repo:**

- templates/tools-pipeline.yaml

#### **Benefit:**

Reduces compliance burden, eliminates third-party tools, and ensures **ATO alignment**.

### **4.3 Tools Inside CodeBuild (Syft, Trivy, Cosign)**

#### **Understanding:**

Pipeline security depends on **artifact transparency and cryptographic assurance**.

#### **Approach:**

- **Syft:** Generates SBOMs in CycloneDX format.
- **Trivy:** Scans for vulnerabilities, failing builds on HIGH/CRITICAL.
- **Cosign:** Signs images + SBOMs with AWS KMS asymmetric key; attestations stored alongside images.

#### **Proof in Repo:**

- ops/pipeline/buildspec-build.yml

**Benefit:**

Ensures **supply chain integrity**, satisfies EO 14028, NIST 800-53, and FedRAMP vulnerability management.

**4.4 FIPS Enforcement (VPC Endpoints)****Understanding:**

FIPS 140-2 requires validated cryptographic modules for all service calls.

**Approach:**

- CodeBuild jobs run inside private subnets, no internet egress.
- Required VPC endpoints validated via `deploy/preflight-vpc.sh`.
- All API calls routed through FIPS DNS suffixes for S3, ECR, KMS, STS, CodeBuild, CodeCommit.

**Proof in Repo:**

- `deploy/preflight-vpc.sh`
- `params/tools-*.json`

**Benefit:**

Guarantees **federal-grade compliance** while preventing data exfiltration.

**4.5 Repo Proof: Buildspecs + Pipeline Templates****Understanding:**

Auditors need **traceability from requirements** → **repo artifacts**.

**Approach:**

- **Infrastructure:** Defined in `templates/tools-pipeline.yaml`.
- **Buildspecs:**
  - `ops/pipeline/buildspec-build.yml` → Build, scan, sign, attest
  - `ops/pipeline/buildspec-deploy.yml` → Deploy to EKS clusters by digest
  - `ops/pipeline/buildspec-mirror-gov.yml` → Mirror signed artifacts into GovCloud
- **IAM Roles:** Defined in `templates/workload-deploy-iam.yaml`

**Benefit:**

Delivers **deterministic, auditable pipelines**, with compliance evidence stored in S3 (SBOMs, vulnerability reports, attestations).

**Chapter 5 — Security & Compliance****5.1 Runtime Enforcement (Kyverno verifyImages, PodSecurity)****Understanding:**

Section 9.2 of the architecture requires **runtime enforcement of supply chain integrity and baseline pod security**.

**Approach:**

- **verifyImages Policy:** Kyverno admission controller ensures only container images signed with KMS-backed Cosign keys are admitted.

- **PodSecurity Policies:** Enforce non-root users, read-only root filesystem, seccomp profiles, and dropped Linux capabilities.
- **Fail Fast:** Unsigned or non-compliant pods are blocked before runtime.

**Proof in Repo:**

- ops/policies/kyverno-verifyimages.yaml
- ops/policies/kyverno-podsecurity.yaml

**Benefit:**

Guarantees **Zero Trust at runtime**, provides auditors with **evidence of enforcement**, and blocks risky workloads.

## 5.2 IAM Least Privilege

**Understanding:**

Strict IAM role separation prevents privilege escalation and aligns with NIST AC family controls.

**Approach:**

- **Environment-Specific Roles:** Separate deploy roles for dev, stg, prod.
- **Scoped Permissions:** Deploy roles grant only EKS + Helm privileges.
- **No Long-Lived Keys:** All role assumption done via IRSA/OIDC federation.

**Proof in Repo:**

- templates/workload-deploy-iam.yaml

**Benefit:**

Minimizes blast radius, enforces least privilege, and aligns with **FedRAMP/NIST 800-53 AC controls**.

## 5.3 IRSA ServiceAccounts & OIDC Federation

**Understanding:**

Static credentials are not compliant in a multi-tenant SaaS. Workloads must assume scoped IAM roles dynamically.

**Approach:**

- **Cluster OIDC Provider:** Each EKS cluster is integrated with OIDC.
- **IRSA:** ServiceAccounts in Kubernetes map to scoped IAM roles.
- **Tenant Workloads:** Each tenant namespace uses its own ServiceAccount + IAM role.

**Proof in Repo:**

- templates/workload-deploy-iam.yaml
- OIDC configuration in deploy/bootstrap-fill-and-deploy.sh

**Benefit:**

Eliminates hardcoded secrets, enables **per-tenant IAM isolation**, and supports CAC/PIV SAML federation.

## 5.4 Zero Trust Posture

### Understanding:

No workload, user, or service is implicitly trusted. Continuous verification is required across the stack.

### Approach:

- **Network:** Private-only subnets + VPC endpoints enforce FIPS traffic.
- **Identity:** OIDC federation + IAM least privilege.
- **Runtime:** Kyverno enforces signed artifacts and pod security baselines.
- **Pipeline:** CodePipeline approvals provide human-in-the-loop governance.

### Proof in Repo:

- `deploy/preflight-vpc.sh`
- `ops/policies/*`
- `templates/tools-pipeline.yaml`

### Benefit:

Aligns with **OMB M-22-09 Zero Trust strategy**, ensures constant verification, and prevents lateral compromise.

## 5.5 Compliance Tie-In (FedRAMP, RMF, FIPS baked in)

### Understanding:

Compliance is baked into the repo — not bolted on.

### Approach:

- **FedRAMP High:** Only authorized AWS services (EKS, CodePipeline, CodeBuild, S3, ECR).
- **FIPS 140-2:** All service calls route through validated endpoints.
- **RMF (NIST 800-37):** Controls mapped across identity, access, monitoring, backup.
- **Evidence:** SBOMs, vulnerability scans, cosign attestations stored in encrypted S3.

### Proof in Repo:

- `ops/pipeline/buildspec-build.yml`
- `ops/policies/kyverno-verifyimages.yaml`
- `templates/tools-pipeline.yaml`

### Benefit:

Provides a **compliance-first platform**, reducing ATO timelines and delivering audit evidence with every build.

## 5.6 Secrets Management (Secrets Manager + External Secrets Operator)

### Understanding:

Applications require access to DB creds, API tokens, and SMTP keys. Hardcoding or storing in ConfigMaps violates compliance.

### Approach:

- **AWS Secrets Manager / SSM Parameter Store:** Stores sensitive values with KMS encryption.

- **External Secrets Operator (ESO):** Syncs secrets from AWS into Kubernetes namespaces dynamically.
- **IAM Controls:** IRSA restricts access so only the right tenant workloads pull their secrets.

**Proof in Repo:**

- ops/policies/app-iam.yaml (workload access policies)
- ESO references in deploy/bootstrap-fill-and-deploy.sh

**Benefit:**

Eliminates secret sprawl, provides rotation capabilities, and ensures **audit logs for every secret access**.

## Chapter 6 — Supply Chain Security

### 6.1 SBOM Generation (Syft)

**Understanding:**

Executive Order 14028 and NIST controls require full transparency into software supply chains via **Software Bill of Materials (SBOMs)**.

**Approach:**

- Every build generates a CycloneDX-format SBOM using **Syft**.
- SBOM artifacts stored in S3 alongside logs and images.
- SBOMs also published to OCI registries as attestations.

**Tools & Platforms:**

- Syft (via CodeBuild container)
- AWS CodeBuild
- Amazon S3

**Proof in Repo:**

- ops/pipeline/buildspec-build.yml (SBOM generation step)

**Benefit:**

Provides visibility into dependencies, accelerates vulnerability management, and satisfies **FedRAMP/NIST SBOM mandates**.

### 6.2 Vulnerability Scanning (Trivy)

**Understanding:**

FedRAMP requires continuous monitoring and vulnerability management. Builds must fail if HIGH/CRITICAL vulnerabilities exist.

**Approach:**

- Container images scanned in CodeBuild using **Trivy** against CVE databases.
- Policy: fail build if HIGH/CRITICAL issues are detected.
- Scan reports archived in encrypted S3.

**Tools & Platforms:**

- Trivy
- AWS CodeBuild
- Amazon S3

**Proof in Repo:**

- ops/pipeline/buildspec-build.yml (scan stage)

**Benefit:**

Prevents insecure workloads from reaching runtime, aligns with **NIST 800-53 CM-8** (System Component Inventory).

**6.3 Signing + Attestation with KMS (Cosign)****Understanding:**

Unsigned artifacts can be tampered with in transit. Compliance requires **cryptographic verification** of builds.

**Approach:**

- **Cosign** signs container images and SBOMs using **AWS KMS asymmetric keys**.
- Signatures stored in ECR alongside images.
- Attestations prove the image was built, scanned, and signed in a controlled environment.

**Tools & Platforms:**

- Cosign
- AWS KMS (asymmetric keys)
- Amazon ECR

**Proof in Repo:**

- ops/pipeline/buildspec-build.yml (sign + attest step)

**Benefit:**

Guarantees artifact provenance, integrity, and compliance with **NIST 800-53 SC-12** and FedRAMP requirements.

**6.4 Kyverno Admission Checks (verifyImages, PodSecurity)****Understanding:**

Even with signatures, enforcement must occur at runtime to prevent unsigned workloads from bypassing the pipeline.

**Approach:**

- Kyverno admission controller enforces **verifyImages** — rejecting pods that lack valid cosign signatures.
- Additional PodSecurity rules enforce non-root, seccomp, and read-only root FS.
- Policies applied cluster-wide at bootstrap.

**Tools & Platforms:**

- Kyverno Admission Controller

**Proof in Repo:**

- ops/policies/kyverno-verifyimages.yaml
- ops/policies/kyverno-podsecurity.yaml

**Benefit:**

Ensures **Zero Trust for workloads**, blocks unsigned or tampered containers, and provides runtime compliance proof.



## 6.5 Immutable ECR Repositories

### Understanding:

Images must be immutable to prevent post-build tampering.

### Approach:

- ECR repos configured as immutable — tags cannot be overwritten.
- Lifecycle policies remove stale images, reducing attack surface.
- Scan-on-push enabled for early vulnerability detection.

### Tools & Platforms:

- Amazon ECR
- IaC via CloudFormation

### Proof in Repo:

- templates/tools-pipeline.yaml (ECR config)

### Benefit:

Guarantees a **tamper-proof chain of custody** from build → deploy, supporting audit requirements for supply chain integrity.

## Chapter 7 — Tenant Lifecycle

### 7.1 Namespace-per-Tenant

#### Understanding:

To support **multi-tenancy** without data leakage, each tenant must have its own namespace boundary inside EKS.

#### Approach:

- A namespace is provisioned for each tenant.
- Tenant-specific objects (ConfigMaps, Secrets, Deployments, Services) live inside the namespace.
- Labels and annotations applied for governance, billing, and observability.

#### Tools & Platforms:

- Amazon EKS
- Kubernetes Namespaces
- Helm templates

#### Proof in Repo:

- ops/tenant/values-tenant-template.yaml

#### Benefit:

Provides **logical isolation**, supports **tenant tagging for billing**, and meets FedRAMP multi-tenancy requirements.

### 7.2 Helm Tenant Templates

#### Understanding:

Onboarding must be **repeatable and automated**. Helm templates standardize tenant configuration.

**Approach:**

- values-tenant-template.yaml defines tenant defaults (quotas, ingress, policies).
- Operators fill in tenant-specific values (TenantID, quotas, RBAC).
- Helm deploys workloads consistently across dev, stg, prod.

**Tools & Platforms:**

- Helm
- Kubernetes
- Amazon EKS

**Proof in Repo:**

- ops/tenant/values-tenant-template.yaml

**Benefit:**

Reduces onboarding time, eliminates misconfigurations, and ensures **consistency across environments**.

### 7.3 Quotas & Policies

**Understanding:**

Without guardrails, tenants could overconsume resources or impact neighbors (“noisy neighbor problem”).

**Approach:**

- **ResourceQuotas:** Cap CPU, memory, and storage per tenant.
- **LimitRanges:** Enforce per-container request/limit ranges.
- **NetworkPolicies:** Restrict pod-to-pod and pod-to-service traffic, preventing lateral movement.
- **RBAC:** Scopes tenant admin roles to their namespace only.

**Tools & Platforms:**

- Kubernetes ResourceQuotas, LimitRanges, NetworkPolicies, RBAC

**Proof in Repo:**

- ops/k8s/tenant-resourcequota.yaml
- ops/k8s/tenant-networkpolicy.yaml

**Benefit:**

Prevents noisy-neighbor risk, enforces **data and traffic isolation**, and satisfies Zero Trust requirements.

### 7.4 Onboarding Automation (bootstrap-fill-and-deploy.sh)

**Understanding:**

Onboarding must be **fast, auditable, and low-risk**. Manual steps introduce errors.

**Approach:**

- bootstrap-fill-and-deploy.sh provisions namespaces, applies Helm tenant templates, and injects IAM roles.
- Guardrails (NetworkPolicies, ResourceQuotas) applied automatically at Day 1.
- Outputs (Helm values, IAM ARNs, policies) logged to S3 for audit evidence.

**Tools & Platforms:**

- Shell automation
- Helm
- IAM + IRSA

**Proof in Repo:**

- `deploy/bootstrap-fill-and-deploy.sh`

**Benefit:**

Enables **rapid, repeatable tenant onboarding** while ensuring compliance guardrails are applied immediately.

**7.5 Noisy-Neighbor Protection & Isolation Models****Understanding:**

The SaaS must serve both **commercial tenants** (cost-optimized pooled/siloed) and **federal tenants** (isolated dedicated stamps).

**Approach:**

- **Pooled/Silo Models:** Use quotas, network policies, and RBAC to enforce logical isolation.
- **Dedicated Stamps:** Provision full stacks (VPC, EKS, CI/CD) for agency workloads requiring maximum assurance.
- **Noisy Neighbor Protection:** LimitRanges + Quotas enforce fairness in pooled models.

**Proof in Repo:**

- `ops/tenant/values-tenant-template.yaml`
- `ops/k8s/tenant-networkpolicy.yaml`
- `ops/k8s/tenant-resourcequota.yaml`

**Benefit:**

Provides **mission-appropriate isolation**, prevents tenant conflicts, and supports both **commercial efficiency** and **federal compliance requirements**.

**Chapter 8 — Public vs Private Partition (Commercial vs GovCloud)****8.1 Commercial → GovCloud Mirroring****Understanding:**

GovCloud has a **restricted AWS service catalog** and cannot run advanced build/scanning stages. To remain compliant, all artifacts must originate from a trusted pipeline in Commercial, then be mirrored into GovCloud.

**Approach:**

- **Builds in Commercial:** CodePipeline + CodeBuild generate SBOMs, run Trivy scans, and sign artifacts with Cosign/KMS.
- **Mirroring:** Signed artifacts mirrored into GovCloud ECR via `gov-mirror-receiver.yaml`.
- **Trust Boundary:** GovCloud workloads only consume artifacts that are **already built, scanned, signed, and attested**.

**Tools & Platforms:**

- AWS CodePipeline (Commercial)

- AWS ECR cross-partition replication
- AWS KMS signing (Cosign integration)

**Proof in Repo:**

- ops/pipeline/buildspec-mirror-gov.yml
- templates/gov-mirror-receiver.yaml

**Benefit:**

Maintains a **single trusted build pipeline**, enforces artifact provenance, and prevents duplicate infrastructure in GovCloud.

## 8.2 GovCloud Restrictions (No Public Endpoints)

**Understanding:**

GovCloud mandates **FIPS endpoints and private-only execution**. Public endpoints and internet egress are prohibited.

**Approach:**

- All GovCloud builds run in private subnets.
- Required VPC endpoints created for ECR, S3, STS, KMS, Logs, CodeBuild.
- **No NAT Gateway** configured — workloads cannot reach public internet.
- Mirrored artifacts from Commercial provide the only ingress into GovCloud pipelines.

**Tools & Platforms:**

- AWS GovCloud (US) partition
- VPC endpoints + private subnets
- CodeBuild (VPC-only mode)

**Proof in Repo:**

- deploy/preflight-vpc.sh (validates endpoints)
- templates/tools-pipeline.yaml

**Benefit:**

Guarantees **FIPS 140-2 compliance**, prevents data exfiltration, and meets ITAR/CJIS/DoD SRG restrictions.

## 8.3 Why This Dual Setup is Necessary

**Understanding:**

Federal workloads cannot operate solely in AWS Commercial due to regulatory constraints. At the same time, GovCloud lacks the service richness needed for SBOMs, scanning, and signing. A dual-partition design satisfies both.

**Approach:**

- **Commercial:** Full build + scan + sign with SBOM generation.
- **GovCloud:** Restricted runtime environment that consumes only pre-validated artifacts.
- **Parity:** Same Helm + IaC deployed in both partitions for consistency.

**Benefit:**

- Meets **federal compliance requirements** (ITAR, DoD SRG, CJIS).
- Leverages **Commercial flexibility** for advanced CI/CD.

- Provides **audit-ready consistency** between Commercial and GovCloud.

## Chapter 9 — Approval & Governance

### 9.1 Manual Approval Stages (dev → stg → prod)

#### Understanding:

In regulated SaaS environments, **human-in-the-loop approvals** are mandatory before workloads advance into higher environments.

#### Approach:

- Manual approval actions inserted between **dev** → **stg** and **stg** → **prod** in CodePipeline.
- Approvers restricted by IAM RBAC (e.g., release managers, security officers).
- Approvals require review of SBOMs, Trivy reports, and cosign attestations before release.

#### Tools & Platforms:

- AWS CodePipeline Manual Approval
- AWS IAM

#### Proof in Repo:

- templates/tools-pipeline.yaml

#### Benefit:

Adds a **compliance gate** that prevents unverified artifacts from being promoted into production.

### 9.2 Signed Artifact Verification

#### Understanding:

Auditors require cryptographic proof that only **signed, attested artifacts** are deployed.

#### Approach:

- **Pipeline Gate:** CodePipeline buildspecs validate cosign signatures before Helm deploy.
- **Runtime Admission:** Kyverno verifyImages policy blocks unsigned workloads from entering the cluster.
- **Dual Enforcement:** Both pipeline and runtime checks required.

#### Tools & Platforms:

- Cosign + AWS KMS
- Kyverno Admission Controller

#### Proof in Repo:

- ops/pipeline/buildspec-deploy.yml
- ops/policies/kyverno-verifyimages.yaml

#### Benefit:

Provides **end-to-end provenance**, from pipeline to runtime, ensuring **supply chain integrity**.

### 9.3 Notifications & Event Routing (SNS, EventBridge, Slack/ServiceNow)

#### Understanding:

Governance requires **visibility into pipeline events** and approvals for stakeholders.

#### Approach:

- **SNS Topics:** Send notifications on build success/failure, approval needed, and deploy complete.
- **EventBridge Rules:** Capture pipeline state changes, enabling downstream workflows (e.g., ServiceNow ticket creation, Slack alerts).
- **Email / ChatOps Integration:** Optional subscription for direct stakeholder alerts.

#### **Tools & Platforms:**

- Amazon SNS
- Amazon EventBridge
- Third-party integrations (Slack, ServiceNow)

#### **Proof in Repo:**

- templates/pipeline-notifications.yaml

#### **Benefit:**

Provides **real-time visibility** into pipeline events, ensures accountability, and maintains **audit trails**.

## **9.4 Ops Handoff**

#### **Understanding:**

Once workloads hit production, **operations teams** assume responsibility for monitoring, patching, and incident response.

#### **Approach:**

- Signed artifact metadata (SBOMs, signatures, vulnerability reports) delivered to Ops via S3 + SNS.
- Ops dashboards (Grafana, CloudWatch) provide observability and cost overlays.
- CI/CD pipeline ensures all evidence is preserved for compliance audits.

#### **Tools & Platforms:**

- Amazon S3
- SNS handoff notifications
- CloudWatch Dashboards

#### **Proof in Repo:**

- ops/pipeline/buildspec-deploy.yml
- templates/pipeline-notifications.yaml

#### **Benefit:**

Ensures a **clear separation of responsibilities**, provides **immutable compliance evidence** for Ops, and reduces deployment risk.

## **Chapter 10 — Observability & Operations**

### **10.1 Logging (CloudWatch, S3, Fluent Bit, ELK optional)**

#### **Understanding:**

Compliance requires **centralized, immutable log storage** with retention policies. Logs must be queryable for audits and incident response.

#### **Approach:**

- **CloudWatch Logs:** EKS cluster logs, application logs, and pipeline logs ingested by default.
- **S3 Archival:** Logs shipped to encrypted S3 with lifecycle policies for long-term retention.
- **Fluent Bit/ELK (Optional):** Forward logs to Elasticsearch for advanced analysis if required.

#### **Tools & Platforms:**

- Amazon CloudWatch Logs
- Amazon S3
- Fluent Bit / ELK (optional)

#### **Proof in Repo:**

- ops/helm/templates/fluent-bit.yaml

#### **Benefit:**

Provides **immutable audit trails**, supports RMF evidence requirements, and enables forensic investigation.

### **10.2 Metrics (CloudWatch Container Insights, AMP/Grafana)**

#### **Understanding:**

Metrics provide visibility into tenant workloads, cluster health, and SLA compliance.

#### **Approach:**

- **CloudWatch Container Insights:** Captures CPU, memory, disk, and network metrics.
- **Amazon Managed Prometheus (AMP):** Scrapes Prometheus exporters at scale.
- **Amazon Managed Grafana (AMG):** Provides dashboards with per-tenant and system-wide visibility.

#### **Tools & Platforms:**

- CloudWatch Container Insights
- AMP + AMG
- Grafana dashboards (dashboards/\*.json)

#### **Proof in Repo:**

- ops/helm/templates/prometheus-exporter.yaml
- dashboards/\*

#### **Benefit:**

Supports **proactive monitoring**, SLA validation, and cost/performance optimization.

### **10.3 Tracing (X-Ray / OpenTelemetry)**

#### **Understanding:**

Multi-tenant SaaS requires tracing across services for troubleshooting and compliance monitoring.

#### **Approach:**

- **AWS X-Ray:** Provides request tracing across APIs and microservices.
- **OpenTelemetry (OTel):** Standardizes traces, integrates with X-Ray and external observability platforms.
- **Tenant-Aware Tags:** Traces labeled with TenantID for per-customer visibility.

**Tools & Platforms:**

- AWS X-Ray
- OpenTelemetry SDK + Collector

**Proof in Repo:**

- ops/k8s/otelsidecar.yaml

**Benefit:**

Improves **MTTD/MTTR**, enables tenant-specific performance reporting, and supports compliance investigations.

**10.4 Cost Allocation (Tags, CUR, Grafana Overlays)****Understanding:**

Accurate cost allocation per tenant is required for billing, chargeback, and compliance reporting.

**Approach:**

- **Resource Tagging:** TenantID, Environment, CostCenter applied to all resources.
- **AWS CUR (Cost & Usage Report):** Filtered by tags to generate per-tenant breakdowns.
- **Grafana Dashboards:** Display cost overlays alongside performance metrics.

**Tools & Platforms:**

- AWS Resource Tagging
- AWS CUR + Cost Explorer
- Grafana

**Proof in Repo:**

- templates/tools-pipeline.yaml (tag enforcement)
- ops/tenant/values-tenant-template.yaml

**Benefit:**

Provides **transparent cost reporting**, ensures fairness, and enables compliance-friendly billing models.

**10.5 Backup, DR & Lifecycle Management****Understanding:**

Disaster Recovery (DR) is required under **FedRAMP contingency planning controls**.

**Approach:**

- **EBS/EFS Backups:** Scheduled backups via AWS Backup, encrypted with KMS.
- **Cross-Region Replication:** Critical S3 buckets and ECR images replicated for redundancy.
- **Velero:** Backups Kubernetes cluster state and tenant namespaces.
- **Lifecycle Policies:** Automatically expire old images and logs.

**Tools & Platforms:**

- AWS Backup
- Velero
- Amazon S3 / ECR lifecycle policies

**Proof in Repo:**



- ops/k8s/backup-policies.yaml
- templates/tools-pipeline.yaml

**Benefit:**

Ensures **continuity of operations**, validates RTO/RPO objectives, and delivers audit-ready recovery evidence.

## 10.6 Cost & Resilience Guardrails

**Understanding:**

Multi-tenant SaaS must balance **efficiency and resilience** to avoid unnecessary spend.

**Approach:**

- **AWS Budgets:** Monitors tenant-level costs and triggers alerts.
- **ECR Lifecycle Policies:** Prevent registry bloat by retaining ~100 images.
- **Resilience Drills:** Regular failover testing for DR validation.
- **Karpenter Replacement (if GovCloud compliant):** For Commercial, Karpenter consolidates workloads; in GovCloud, use managed node groups.

**Tools & Platforms:**

- AWS Budgets
- AWS Backup + DR runbooks
- ECR lifecycle management

**Proof in Repo:**

- ops/k8s/backup-policies.yaml
- templates/tools-pipeline.yaml

**Benefit:**

Delivers **predictable costs**, enforces **tenant fairness**, and strengthens **resilience posture** across environments.

## Chapter 11 — Criteria Compliance Matrix

**Understanding:**

Section 18 of the AWS\_MultiTenant\_SaaS\_EKS\_FIPS\_Architecture requires a **repeatable, auditable compliance mapping**. This matrix shows how repo artifacts fulfill FedRAMP, FIPS 140-2, NIST 800-53, and RMF controls.

Control Requirement	Architecture Doc Section	Repo File(s)	Compliance Justification
<b>Supply Chain Integrity</b>	Sec 9.2 (Platform Security Controls)	ops/pipeline/buildspec-build.yml, ops/policies/kyverno-verifyimages.yml	All images scanned, signed, attested, and enforced at runtime. Ensures EO 14028 + FedRAMP compliance.
<b>Tenant Isolation</b>	Sec 4 (Multi-Tenant Isolation Models)	ops/tenant/values-tenant-template.yml, ops/k8s/tenant-networkpolicy.yml, ops/k8s/tenant-resourcequota.yml	Enforces namespace, network, and resource boundaries. Prevents cross-tenant access.
<b>FIPS Endpoint Enforcement</b>	Sec 3.3 & 10.2	deploy/preflight-vpc.sh, templates/tools-pipeline.yml, params/tools-*.json	All builds + API calls restricted to private subnets and FIPS 140-2 validated endpoints.
<b>Least Privilege IAM</b>	Sec 5.2 (IAM Guardrails)	templates/workload-deploy-iam.yml	Deploy roles scoped per environment; workloads assume scoped IRSA roles.
<b>Runtime Security</b>	Sec 9.2	ops/policies/kyverno-podsecurity.yml	Pods run non-root, with seccomp, read-only FS, and dropped capabilities.
<b>CI/CD Pipeline Security</b>	Sec 10 (CI/CD Requirements)	ops/pipeline/buildspec-deploy.yml, ops/pipeline/buildspec-mirror-gov.yml, templates/tools-pipeline.yml	Pipeline enforces immutable digests, manual approvals, and GovCloud mirroring.
<b>Identity Federation</b>	Sec 10.2 (Identity & Access)	templates/workload-deploy-iam.yml, OIDC configs	Cognito/OIDC federation for users; IRSA for workloads; no static credentials.

Control Requirement	Architecture Doc Section	Repo File(s)	Compliance Justification
<b>Approvals &amp; Governance</b>	Sec 10 (Approvals)	templates/tools-pipeline.yaml, templates/pipeline-notifications.yaml	Manual approvals inserted in pipeline; SNS/EventBridge notify stakeholders.
<b>Observability</b>	Sec 12 (Operational Guardrails)	ops/helm/templates/fluent-bit.yaml, ops/k8s/otel-sidecar.yaml, dashboards/*	Logs, metrics, and traces captured via CW/AMP/AMG/X-Ray.
<b>Backup &amp; DR</b>	Sec 18 (Implementation Plan)	ops/k8s/backup-policies.yaml, templates/tools-pipeline.yaml	Automated snapshots, cross-region replication, and recovery drills ensure continuity.
<b>Immutable Repos</b>	Sec 10.2 (Artifact Control)	templates/tools-pipeline.yaml (ECR config)	ECR repos immutable, lifecycle policies enabled, prevents overwrite tampering.
<b>Compliance Evidence</b>	Sec 18 (Audit Readiness)	S3 artifact storage, ops/pipeline/buildspec-build.yml	SBOMs, scan reports, signatures archived securely for auditors.
<b>Secrets Management</b>	Sec 9.2 (Security Controls)	ops/policies/app-iam.yaml, deploy/bootstrap-fill-and-deploy.sh	Secrets in AWS Secrets Manager, synced with External Secrets Operator; per-tenant IRSA access.
<b>Zero Trust Enforcement</b>	OMB M-22-09, Sec 9.2	ops/policies/kyverno-verify-images.yaml, deploy/preflight-vpc.sh	Continuous verification enforced across network, identity, runtime, and pipeline.

#### Benefit:

This matrix makes the repo an **audit-ready package**: every requirement maps to a file and justification, allowing compliance teams to validate FedRAMP, FIPS, and RMF quickly.

## Chapter 12 — Conclusion

### 12.1 How the Repo + Process Fulfill Every Requirement

The **AWS EKS CI/CD (Commercial + GovCloud) via CloudFormation** repository is not just infrastructure code — it is a **compliance-enforcing SaaS delivery framework**. Every repo artifact maps directly to architecture requirements and federal mandates.

- **Multi-Tenant SaaS Principles (Section 4):**  
Implemented via Kubernetes namespaces, Helm templates, ResourceQuotas, and NetworkPolicies (ops/tenant/\*, ops/k8s/\*). Supports pooled, siloed, and dedicated tenant models from a single codebase.
- **Infrastructure and FIPS Controls (Sections 3.3 & 10.2):**  
VPC-only subnets, private endpoints, and FIPS-validated service calls enforced via preflight validation (deploy/preflight-vpc.sh, templates/tools-pipeline.yaml).
- **CI/CD Pipeline Security (Section 10):**  
End-to-end CodePipeline stages enforce build → scan → sign → attest → approval → deploy → mirror. Repo-defined buildspecs (ops/pipeline/\*) and IAM templates (templates/workload-deploy-iam.yaml) ensure controlled promotion and GovCloud parity.
- **Supply Chain Security (Section 9.2):**  
SBOM generation (Syft), vulnerability scanning (Trivy), signing + attestations (Cosign), and runtime enforcement (Kyverno policies) guarantee provenance and integrity.
- **Observability and Ops Guardrails (Section 12):**  
Logging (CloudWatch/Fluent Bit), metrics (AMP/AMG), tracing (X-Ray/OTel), and dashboards (dashboards/\*) deliver full tenant-aware visibility and SLA monitoring.
- **Compliance Evidence and RMF Tie-In (Section 18):**  
All SBOMs, scan reports, signatures, and pipeline logs stored in encrypted S3, producing an immutable audit trail for FedRAMP, FIPS, and RMF validation.

### 12.2 Why This Design Is Production-Ready, Auditable, and Extensible

#### Production-Ready:

- Runs on AWS-native FedRAMP High services (EKS, CodePipeline, CodeBuild, GovCloud).
- Enforces Zero Trust across network, identity, runtime, and CI/CD layers.
- Proven onboarding automation (deploy/bootstrap-fill-and-deploy.sh) for tenants.

#### Auditable:

- Cryptographic signatures and SBOMs provide **tamper-proof provenance**.
- Runtime policies enforce compliance continuously, not just at build time.
- Criteria Compliance Matrix (Chapter 11) provides **direct audit mapping**.

#### Extensible:

- Single repo supports pooled, siloed, and dedicated stamps.
- Helm + CloudFormation modularity allows adding new tenants, policies, and services without architectural redesign.

- Dual-partition (Commercial + GovCloud) model ensures consistent operations across sensitive and regulated environments.

## Final Statement

This repository and proposal provide a **blueprint SaaS platform** that is:

- **Secure:** Zero Trust enforced across compute, identity, and runtime.
- **Compliant:** FedRAMP, FIPS 140-2, and RMF controls built into every layer.
- **Scalable:** Supports elastic onboarding across pooled, siloed, and dedicated models.
- **Auditable:** Immutable evidence stored in encrypted S3, traceable to repo files.
- **Extensible:** Modular architecture ensures future-proof growth and customer adoption.

It is a **production-ready SaaS framework** for Commercial and Government customers, designed to withstand compliance scrutiny, deliver operational excellence, and enable rapid, secure adoption of Zero Trust principles.

## Appendices

### Appendix A — AWS Services by Function

#### 1. SaaS Platform / Application Layer

- **Amazon EKS** — Multi-tenant Kubernetes clusters (dev, stg, prod).
- **Amazon ECR** — Immutable container registry (Commercial + GovCloud).
- **Amazon EC2 / Managed Node Groups** — Compute capacity (GovCloud-approved).
- **AWS Fargate (optional)** — Serverless workloads (Commercial only).
- **Amazon VPC** — Networking foundation (multi-AZ private subnets).
- **Amazon Route 53** — DNS + failover routing.
- **AWS Load Balancer Controller (ALB/NLB)** — Ingress for tenant workloads.

#### 2. CI/CD & Supply Chain

- **AWS CodeCommit** — Git repo for source control.
- **AWS CodeBuild** — Builds, scans, and signs workloads.
- **AWS CodePipeline** — Orchestration for build → scan → sign → deploy → mirror.
- **AWS KMS** — Asymmetric keys for Cosign signing.
- **Amazon S3** — Storage for SBOMs, logs, pipeline evidence.
- **Cross-partition ECR** — Artifact mirroring to GovCloud.
- *3rd Party Tools in Build Containers:* Syft (SBOM), Trivy (scanning), Cosign (signing).

#### 3. Security & Compliance

- **AWS IAM** — Role-based access control.
- **IRSA (IAM Roles for Service Accounts)** — Pod-level IAM integration.
- **Amazon Cognito** — End-user authentication with OIDC/SAML federation.
- **AWS WAF** — Web Application Firewall at CloudFront/ALB.
- **AWS CloudTrail** — Account/pipeline activity logging.
- **Kyverno** — Runtime policy enforcement (verifyImages, pod security).
- **AWS PrivateLink (VPC Endpoints)** — FIPS-only API calls.

#### 4. Observability & Operations

- **Amazon CloudWatch Logs** — Centralized logs.
- **CloudWatch Container Insights** — EKS metrics.
- **Amazon Managed Prometheus (AMP)** — Scraping metrics.
- **Amazon Managed Grafana (AMG)** — Dashboards + SLO visibility.
- **AWS X-Ray** — Distributed tracing.
- **OpenTelemetry (OTel)** — Tracing pipeline integration.

#### 5. Governance & Notifications

- **Amazon SNS** — Notifications for approvals/events.
- **Amazon EventBridge** — Routing pipeline state events.
- **Slack/ServiceNow (optional)** — Approval/alert integrations.

#### 6. Data & Backup / DR

- **AWS Backup** — Backup policies for EBS/EFS.
- **Velero** — Cluster/namespace backup.
- **Amazon S3 Cross-Region Replication** — DR redundancy.
- **Amazon DynamoDB / Aurora** — For app data if required.
- **Amazon ElastiCache** — Caching tier.
- **Amazon SQS** — Event decoupling for microservices.

### Appendix B — Environment Setup & Validation Checklist

#### 0) Prereqs

- AWS CLI authenticated to Commercial (and GovCloud if used).
- kubectl + cluster access (eks-dev, eks-stg, eks-prod).
- Helm 3 installed.
- Kyverno installed in all clusters.

#### 1) Config Setup

- Populate deploy/bootstrap.conf (APP\_NAME, REPO\_NAME, VPCs, Subnets, etc.).
- Define Commercial + GovCloud account IDs and regions.
- Provide KMS signing alias (alias/my-service-cosign).
- Configure Cognito or OIDC values.

#### 2) KMS Signing Key

- Create asymmetric ECC key in KMS.
- Alias must match COSIGN\_KMS\_ALIAS.
- Allow codebuild-build-role to sign.

#### 3) Preflight VPC Validation

- Run deploy/preflight-vpc.sh to validate endpoints (ECR, S3, STS, KMS, Logs, CodeBuild).
- Fix missing endpoints before bootstrap.

#### 4) Bootstrap & Deploy

- Run deploy/bootstrap-fill-and-deploy.sh.
- Confirm: params/\*.json, Helm values, tenant templates, stacks created.

## 5) Cluster Prep

- Confirm EKS clusters exist (dev/stg/prod).
- Add-ons: Metrics Server, ALB Controller, OIDC provider.

## 6) Kyverno Setup

- Fetch cosign public key with `deploy/fetch-cosign-pubkey.sh`.
- Apply `verifyImages` + pod security policies (`ops/policies/*`).

## 7) Pipeline Setup

- Push repo to CodeCommit.
- Validate pipeline stages: `build` → `scan` → `sign` → `attest` → `deploy` → `mirror`.

## 8) Ingress & DNS

- Confirm ALB, ACM cert, WAF association, Route53 DNS entry.

## 9) Notifications

- Deploy `templates/pipeline-notifications.yaml` to enable SNS/EventBridge notifications.

## 10) Tenant Onboarding

- Use `ops/tenant/values-tenant-template.yaml`.
- Apply quotas and network policies.
- Deploy tenant workloads by image digest.

## 11) Ops Hardening

- Deploy External Secrets Operator.
- Enable observability (CW, AMP, Grafana, OTel).
- Validate backups (AWS Backup + Velero).

## Appendix C — Deployment Evidence (Scripts & Outputs)

### Scripts

- `deploy/bootstrap-fill-and-deploy.sh` — Automates namespace + tenant onboarding.
- `deploy/preflight-vpc.sh` — Validates FIPS endpoint compliance.
- `deploy/fetch-cosign-pubkey.sh` — Extracts public key for Kyverno `verifyImages` policy.
- `deploy/apply-kyverno.sh` — Applies admission policies.

### Outputs

- **Pipeline Outputs:** Immutable image digests, SBOMs, Trivy reports, cosign attestations.
- **Cluster Evidence:** Namespace manifests, applied quotas, applied Kyverno policies.
- **Compliance Evidence:** S3 storage of all reports/artifacts.

## Appendix D — File-to-Chapter Mapping (Repo Proof)

- **Chapter 3 Infrastructure** → `templates/tools-pipeline.yaml`, `deploy/preflight-vpc.sh`
- **Chapter 4 CI/CD** → `ops/pipeline/buildspec-build.yml`, `ops/pipeline/buildspec-deploy.yml`
- **Chapter 5 Security** → `ops/policies/kyverno-verifyimages.yaml`, `templates/workload-deploy-iam.yaml`

- **Chapter 6 Supply Chain** → ops/pipeline/buildspec-build.yml, templates/tools-pipeline.yml
- **Chapter 7 Tenancy** → ops/tenant/values-tenant-template.yml, ops/k8s/tenant-networkpolicy.yml
- **Chapter 8 Partitioning** → ops/pipeline/buildspec-mirror-gov.yml, templates/gov-mirror-receiver.yml
- **Chapter 9 Governance** → templates/pipeline-notifications.yml
- **Chapter 10 Observability** → ops/k8s/otel-sidecar.yml, ops/k8s/backup-policies.yml, dashboards/\*

#### **Appendix E — Optional AWS Services for Scale-Out (Optional Enhancements)**

- **AWS Control Tower** — Multi-account baseline for governance at scale.
- **AWS Service Catalog** — Self-service tenant onboarding templates.
- **Amazon GuardDuty + Security Hub** — Centralized threat detection.
- **AWS Config + Conformance Packs** — Continuous compliance monitoring.
- **Amazon OpenSearch Service** — Enterprise search/log analytics.